

# Task 3 Classifying Credit Risk

## The box plot

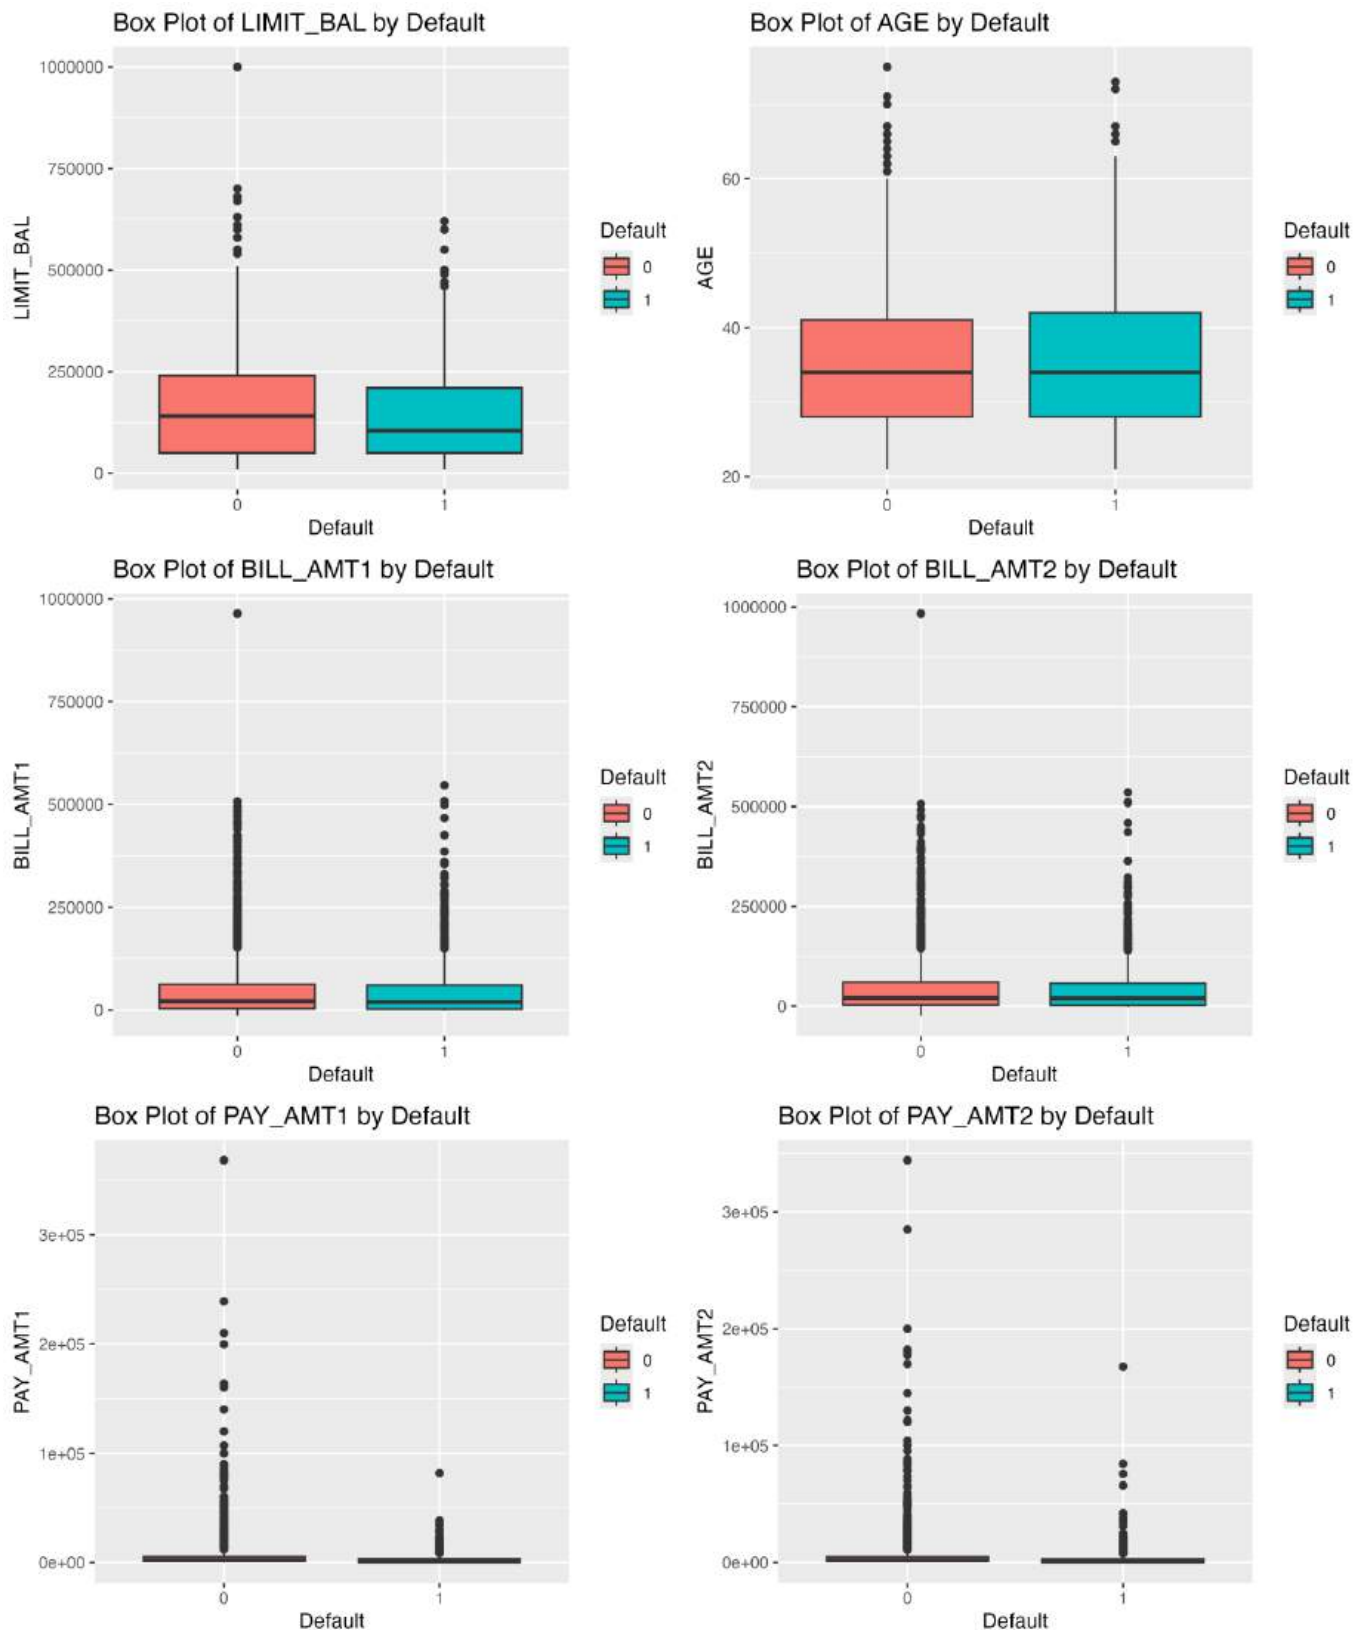


Figure 1: Box Plots for Continuous Variables by Default Status

## Purpose

Box plots are used to visualize the distribution of continuous variables. They provide a summary of key statistics such as the median, quartiles, and potential outliers. In this context, the box plots help compare the distributions of various financial metrics between those who defaulted on their credit card payments (default = 1) and those who did not (default = 0).

## Techniques Used

1. **Data Cleaning:** We used **na.omit(data)** to remove any rows with **NA** values from the dataset. This ensures that our analysis is not skewed by missing data.
2. **Box Plot Creation with ggplot2:**
  - We used the **ggplot2** package to create box plots for continuous variables.
  - The **geom\_boxplot()** function is used to generate the box plot.
  - The **aes()** function specifies the aesthetics, such as the x-axis (default status), y-axis (continuous variable), and fill color (default status).
  - The **labs()** function is used to label the axes and title.
  - The **scale\_fill\_discrete()** function is used to ensure the fill color corresponds to the default values.
3. **Plot Arrangement:**
  - We used the **gridExtra** package to arrange multiple plots in a grid layout.
  - The **grid.arrange()** function is used to arrange the plots in a 2x3 grid.
  - The **ggsave()** function saves the combined plot as an image file.

## Assumptions

1. **Continuity of Data:** The continuous variables (e.g., **LIMIT\_BAL**, **AGE**, **BILL\_AMT1**, etc.) are assumed to be measured on a continuous scale. This allows for the computation of quartiles and the identification of outliers.
2. **Independence of Observations:** Each observation (i.e., each credit card client's data) is assumed to be independent of others. This is important for the validity of the statistical summaries represented in the box plots.
3. **Handling of Outliers:** Box plots inherently show outliers as points beyond the whiskers. It is assumed that these outliers are genuine data points and not errors. Outliers are included to give a complete picture of the data distribution.
4. **Categorical Representation of Default Status:** The default status (**default**) is treated as a categorical variable with two levels (0 and 1). The box plots are filled and separated based on these levels to facilitate comparison.

## The bar chart

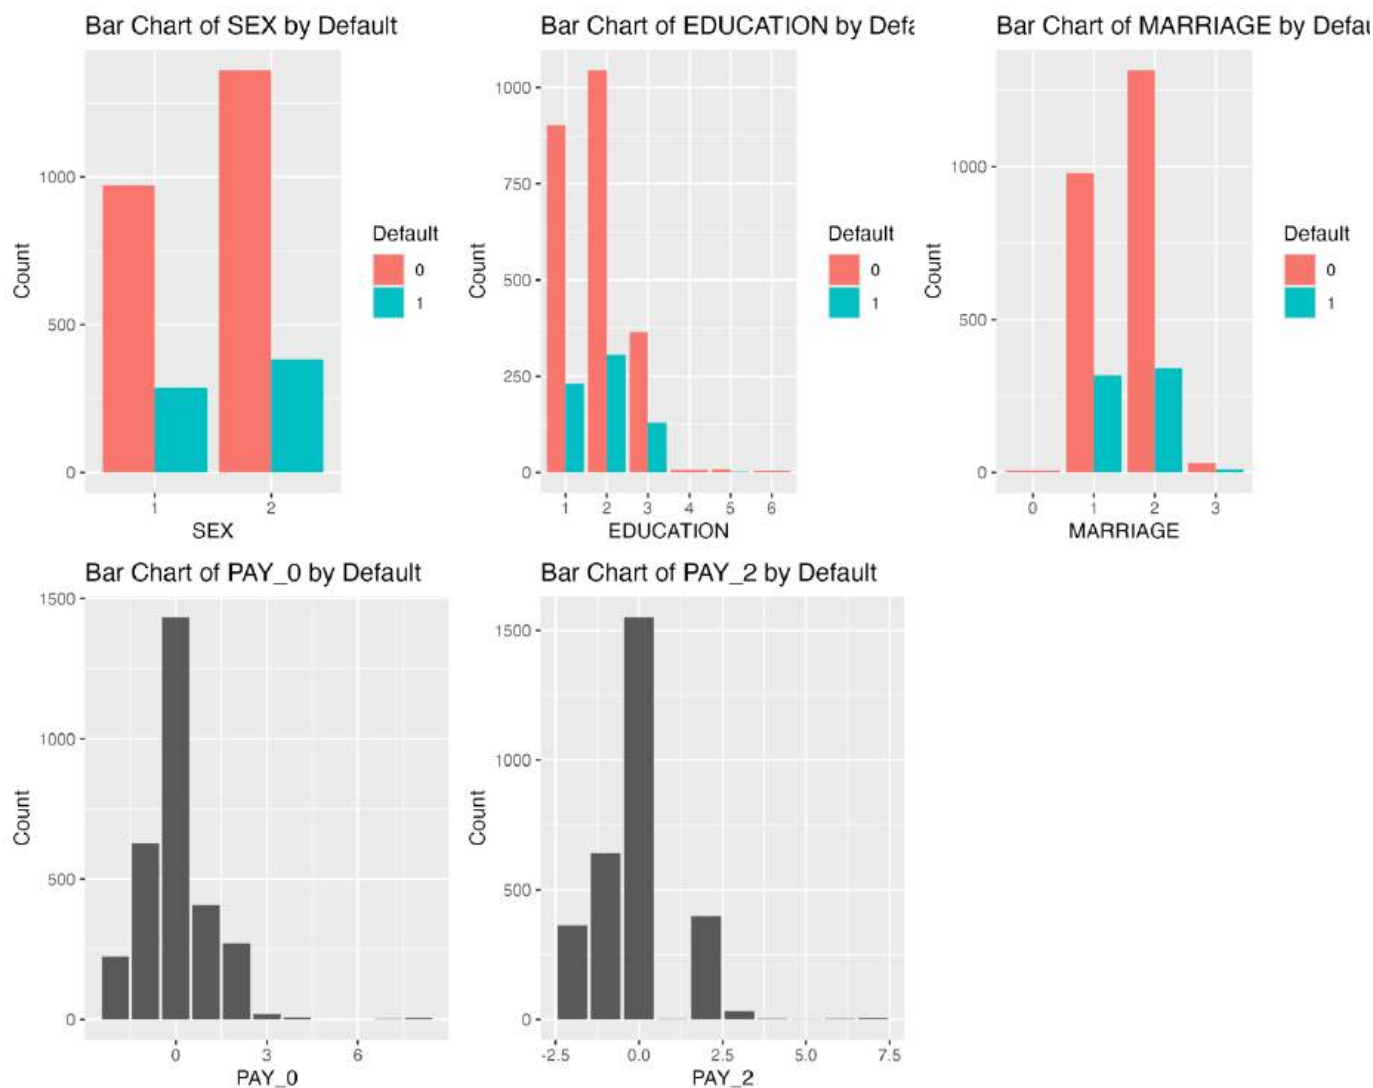


Figure 2: Bar Charts for Categorical Variables by Default Status

## Purpose

Bar charts are used to visualize the distribution of categorical variables. They provide a clear comparison of counts across different categories. In this context, the bar charts help compare the distribution of demographic factors (SEX, EDUCATION, MARRIAGE) and repayment status (PAY\_0, PAY\_2) between those who defaulted on their credit card payments (default = 1) and those who did not (default = 0).

## Techniques Used

1. **Data Cleaning:** Similar to the box plots, `na.omit(data)` is used to remove any rows with **NA** values from the dataset to ensure the analysis is not affected by missing data.
2. **Bar Chart Creation with ggplot2:**
  - We used the **ggplot2** package to create bar charts for categorical variables.
  - The **geom\_bar(position = "dodge")** function is used to generate the bar charts with bars for different default statuses placed side-by-side (dodge position).
  - The **aes()** function specifies the aesthetics, such as the x-axis (categorical variable), fill color (default status), and y-axis (count).
  - The **labs()** function is used to label the axes and title.
  - The **scale\_fill\_discrete()** function is used to ensure the fill color corresponds to the default values.

### 3. Plot Arrangement:

- We used the **gridExtra** package to arrange multiple plots in a grid layout.
- The **grid.arrange()** function is used to arrange the plots in a 2x2 grid.
- The **ggsave()** function saves the combined plot as an image file.

### Assumptions

1. **Categorical Data:** The variables **SEX**, **EDUCATION**, **MARRIAGE**, **PAY\_0**, and **PAY\_2** are treated as categorical variables with distinct levels. This allows for the counting and comparison of occurrences within each category.
2. **Independence of Observations:** Each observation (i.e., each credit card client's data) is assumed to be independent of others. This is important for the validity of the counts represented in the bar charts.
3. **Removed Missing Values:** By using **na.omit(data)**, we assume that the remaining data without **NA** values is representative of the overall dataset. This helps ensure that the analysis is not biased by missing data.

## Discussion on the Variables

### Descriptive Statistics

#### Continuous Variables:

Variable	Min	Mean	Median	Std Dev	Max
LIMIT_BAL	10000	164036.7	140000	128263.9	1,000,000
AGE	21	35.34	34	9.36	75
BILL_AMT1	-14386	50199.78	21253.5	75769.4	964511
BILL_AMT2	-24704	48144.65	20384.5	73539.03	983931
PAY_AMT1	0	5387.124	2091	14354.51	368199
PAY_AMT2	0	5296.843	2000	15661.35	344261

Table 1: Descriptive Statistics for Continuous Variables

#### Categorical Variables:

Variable	Default = 0	Default = 1
SEX_1	970	286
SEX_2	1362	382
EDUCATION_1	902	231
EDUCATION_2	1046	306
EDUCATION_3	364	129

EDUCATION_4	7	0
EDUCATION_5	8	2
EDUCATION_6	5	0
MARRIAGE_0	7	0
MARRIAGE_1	978	318
MARRIAGE_2	1316	340
MARRIAGE_3	31	10
PAY_0_0	1245	189
PAY_0_1	269	139
PAY_0_2	92	178
PAY_0_3	9	11
PAY_0_4	1	7
PAY_0_5	0	0
PAY_0_6	0	0
PAY_0_7	0	2
PAY_0_8	5	1
PAY_0_9	0	0
PAY_2_0	1299	253
PAY_2_1	202	195
PAY_2_2	11	22
PAY_2_3	0	3
PAY_2_4	1	0
PAY_2_5	0	3
PAY_2_6	5	1
PAY_2_7	0	0

<b>PAY_2_8</b>	0	0
<b>PAY_2_9</b>	2	3

Table 2: Descriptive Statistics for Categorical Variables

## Comments on the Variables

### Continuous Variables

#### 1. **LIMIT\_BAL:**

- Both groups (default and non-default) have similar distributions in terms of credit limit.
- The mean and median credit limits are slightly higher for non-defaulters.
- There are outliers in both groups with very high credit limits, indicating some individuals have access to large amounts of credit.

#### 2. **AGE:**

- Age distribution is similar for both groups, with a median age around 34 years.
- The standard deviation indicates a moderate spread around the mean age.
- No significant difference in age between defaulters and non-defaulters.

#### 3. **BILL\_AMT1 and BILL\_AMT2:**

- The distributions of bill amounts for the two months are similar across both groups.
- There are negative values, indicating possible corrections or refunds.
- The mean and median bill amounts are slightly higher for non-defaulters.

#### 4. **PAY\_AMT1 and PAY\_AMT2:**

- Both groups show wide ranges of payment amounts.
- There are high outliers, especially among non-defaulters, indicating some individuals make large payments.
- The median payment amounts are low, showing that most clients make smaller payments.

### Categorical Variables:

#### 1. **SEX:**

- More females (SEX\_2) in the dataset.
- Higher count of defaults among females, but the proportion of defaults is similar for both males and females.

#### 2. **EDUCATION:**

- Most clients have a university (EDUCATION\_2) or graduate school (EDUCATION\_1) education.
- Higher counts of defaults among these education levels, but this aligns with the higher number of individuals in these categories.

#### 3. **MARRIAGE:**

- Majority of clients are single (MARRIAGE\_2) or married (MARRIAGE\_1).
- More defaults among single clients compared to married ones, suggesting marital status might influence default rates.

#### 4. **PAY\_0:**

- Most clients have a PAY\_0 status of 0 (pay duly) for both defaulters and non-defaulters.

- Higher counts of default are observed for higher PAY\_0 values, indicating longer payment delays are associated with higher default rates.

#### 5. PAY\_2:

- Similar to PAY\_0, most clients have a PAY\_2 status of 0 (pay duly).
- Higher counts of default are observed for higher PAY\_2 values, indicating that clients with longer payment delays in August 2005 are more likely to default

### Discussion on Variables

From the graphs and summary statistics, several observations can be made:

- **Credit Limit (LIMIT\_BAL)** and **Age (AGE)** distributions do not show significant differences between the default and non-default groups.
- **Bill Amounts (BILL\_AMT1 and BILL\_AMT2)** have similar distributions across both groups, though higher values are observed among non-defaulters.
- **Payment Amounts (PAY\_AMT1 and PAY\_AMT2)** indicate that non-defaulters tend to make larger payments, reflected in the presence of higher outliers.
- **Gender (SEX)** shows that defaults are more frequent among females, but the proportion of defaults is not significantly different between genders.
- **Education Level (EDUCATION)** reveals that higher education levels (university and graduate school) have higher counts of defaults, possibly due to the larger representation of these groups in the dataset.
- **Marital Status (MARRIAGE)** suggests single individuals are more likely to default compared to married ones.
- **Repayment Status (PAY\_0 and PAY\_2)** indicates that longer delays in payments are associated with higher default rates. Most clients who pay duly (status 0) have lower default rates, while those with higher delays (status 1 and above) show higher default rates. This suggests that recent repayment behavior is a strong indicator of default risk, highlighting the importance of timely payments to avoid default

## Task 4 Forecasting Stock Returns using Technical Indicators

### Outcome Variable:

#### 4.1.1 Daily Logarithmic Returns

- **Description:** Logarithmic returns, or log returns, are calculated as the natural logarithm of the ratio of today's closing price to yesterday's closing price.
- **Formula:**  $\log\_ret = \log\left(\frac{P_{t-1}}{P_t}\right)$
- **Reason for Use:** Logarithmic returns are preferred in financial time series analysis because they are time-additive, making it easier to aggregate returns over multiple periods. They also account for compounding effects.

### Predictor Variables:

#### 4.1.2 One Period Lag of the Log Returns

- **Description:** The log return of the previous day.
- **Reason for Use:** Lagged log returns can capture the momentum effect in stock prices, where past returns might influence future returns.

#### 4.1.3 One Period Lag of Moving Average: 5 day

- **Description:** The simple moving average (SMA) over the past 5 days, lagged by one period.
- **Formula:**  $SMA_5 = \frac{P_{t-1} + P_{t-2} + P_{t-3} + P_{t-4} + P_{t-5}}{5}$
- **Reason for Use:** Moving averages smooth out short-term price fluctuations and help identify trends. The lagged SMA indicates the trend direction from the previous day.

#### 4.1.4 One Period Lag of Exponential Moving Average: 5 day

- **Description:** The exponentially weighted moving average (EMA) over the past 5 days, lagged by one period.
- **Formula:**  $EMA_5 = \left(\frac{2}{n+1}\right) (P_t - EMA_{t-1}) + EMA_{t-1}$  where  $n = 5$
- **Reason for Use:** The EMA gives more weight to recent prices, making it more responsive to new information. The lagged EMA helps in identifying recent trends.

#### 4.1.5 One Period Lag of RSI with 5 Day Period

- **Description:** The Relative Strength Index (RSI) over the past 5 days, lagged by one period.
- **Formula:** RSI is calculated as  $100 - \left(\frac{100}{1+RS}\right)$ , where RS is the average of 'n' days' up closes divided by the average of 'n' days' down closes.
- **Reason for Use:** RSI measures the speed and change of price movements, helping to identify overbought or oversold conditions. The lagged RSI provides insight into past market momentum.

## Sentiment Indicator

#### 4.2 Sentiment Indicator

- **Description:** The daily aggregated sentiment score from Alexandria (ACTA) data, converted to an **xts** time series object.
- **Reason for Use:** Sentiment indicators can capture market emotions and investor behavior, which can influence stock prices. Lagging the sentiment indicator allows us to use past sentiment as a predictor for future stock movements.

## Data Splitting

#### 4.3 Training and Testing Sample:

- **Description:** The dataset is split into training and testing sets, with the last 100 days used for testing and the remaining data for training.
- **Reason for Use:** This split ensures that the model is trained on historical data and tested on recent, unseen data, allowing for an evaluation of its predictive performance.

Date	log_returns	ma_5	ema_5	rsi_5	sentiment
2021-01-12	-0.004280781	90.114	90.24467	57.04324	0.04269946
2021-01-13	-0.018428585	89.826	89.91311	44.24013	0.05108035
2021-01-14	0.005809423	90.180	89.86541	48.74472	0.20270494



2021-01-15	0.002447709	90.244	89.90694	50.84478	0.21189820
2021-01-19	0.007748544	90.122	90.16796	57.73264	0.06954699
2021-01-20	0.026764000	90.570	91.16197	73.83713	0.21663397

Table 3: Sample Training set

## Neural Network Regression Forecasting

### 4.4 Neural Network Regression Forecasting:

- **4.4.1 Training on Training Sample:** Using 'timeslice' sampling with a 100-day horizon and the remaining training days as the initial window.
- **4.4.2 Data Pre-processing:** Standardizing the data to ensure all predictors have a mean of zero and a standard deviation of one.
- **4.4.3 Model Fitting:** Using the `nnet` method with RMSE metric and a `tuneLength` of 5 for model tuning.
- **4.4.4 Prediction on Test Set:**

Metric	Value
RMSE	0.009827654
R-squared	0.001399213
MAE	0.007019770

Table 4: Neural Network results

- **4.4.5 Plot of Predicted and Actual Values:** Visualizing the predictions against actual values to assess model performance.




Figure 3: Neural Network Forecasting: Actual vs Predicted

- 4.4.6 Discussion on Accuracy:** The performance of the Neural Network Regression model for forecasting stock returns is evaluated using three key metrics: RMSE, R-squared, and MAE. The RMSE of 0.0098 indicates that the average deviation of the predicted values from the actual values is quite low, suggesting the model's predictions are relatively accurate. The MAE of 0.0070 further supports this by showing that the average absolute error between the predicted and actual values is minimal.

However, the R-squared value is only 0.0014, indicating that the model explains just 0.14% of the variance in the log returns. This low R-squared suggests that while the model's predictions are close to the actual values in terms of magnitude, it fails to capture the underlying variability in the stock returns. This could be due to the inherent volatility of stock prices, the limitations of the selected technical indicators and sentiment analysis, or the need for further tuning and additional features.

## SVM Regression Forecasting

### 4.5 SVM Regression Forecasting:

- 4.5.1 Radial Kernel:** Implementing SVM regression with a radial basis function (RBF) kernel for capturing non-linear relationships.
- 4.5.2 Prediction on Test Set:**

Metric	Value
RMSE	0.010307150
R-squared	0.005913989
MAE	0.007485382

Table 5: SVM Regression results

- 4.5.3 Plot of Predicted and Actual Values:**




Figure 4: SVM Regression Forecasting: Actual vs Predicted

- 4.5.4 Discussion on Accuracy:** The performance of the SVM Regression model for forecasting stock returns is evaluated using RMSE, R-squared, and MAE. The RMSE of 0.0103 indicates that the average deviation of the predicted values from the actual values is low, suggesting reasonable accuracy. The MAE of 0.0075 further supports this by showing that the average absolute error between the predicted and actual values is minimal. The R-squared value of 0.0059 suggests that the model explains approximately 0.59% of the variance in the log returns, which is still quite low but

slightly better than the Neural Network model.

Comparing the SVM model with the Neural Network model, the Neural Network has a slightly lower RMSE (0.0098) and MAE (0.0070), indicating better predictive accuracy in terms of average error magnitude. However, the R-squared value for the Neural Network is only 0.0014, much lower than the SVM's R-squared of 0.0059, suggesting that the SVM explains more variance in the log returns than the Neural Network.

Based on the predictive performance, both models show similar predictive accuracy, but the Neural Network has a slight edge in terms of RMSE and MAE. However, considering the explanatory power, the SVM is better as it explains more variance. If the goal is purely to minimize prediction error, the Neural Network is recommended. If understanding and explaining the variance in stock returns is also important, the SVM might be more suitable. Overall, given the marginal difference in RMSE and MAE, the Neural Network might be preferred for its slightly better predictive accuracy.

# Reference list

- Loh, W.-Y., 2011. Classification and regression trees. *Wiley Interdisciplinary Reviews: Data Mining and Knowledge Discovery*, 1(1), pp. 14-23. Available at: <https://onlinelibrary.wiley.com/doi/abs/10.1002/widm.8> [Accessed 18 May, 2024].
- Hand, D.J. and Henley, W.E., 1997. Statistical classification methods in consumer credit scoring: a review. *Journal of the Royal Statistical Society: Series A (Statistics in Society)*, 160(3), pp. 523-541. Available at: <https://www.jstor.org/stable/2983089> [Accessed 18 May, 2024].
- Breiman, L., 2001. Random forests. *Machine Learning*, 45(1), pp. 5-32. Available at: <https://link.springer.com/article/10.1023/A:1010933404324> [Accessed 18 May, 2024].
- Quinlan, J.R., 1986. Induction of decision trees. *Machine Learning*, 1(1), pp. 81-106. Available at: <https://link.springer.com/article/10.1007/BF00116251> [Accessed 18 May 2024].
- Jolliffe, I.T., 2002. Principal component analysis. Springer, New York. Available at: <https://link.springer.com/book/10.1007/b98835> [Accessed 18 May, 2024].
- Fan, J., & Yao, Q., 2017. The Elements of Financial Econometrics. Cambridge University Press. ISBN: 978-1107161318. Available at: <https://www.cambridge.org/core/books/elements-of-financial-econometrics/3B20C70EB6F913A2E7C6F8AC1A74D058> [Accessed 18 May 2024].
- Connor, G. and Korajczyk, R.A., 1993. A Test for the Number of Factors in an Approximate Factor Model. *The Journal of Finance*, 48(4), pp.1263-1291. Available at: [https://papers.ssrn.com/sol3/papers.cfm?abstract\\_id=959023](https://papers.ssrn.com/sol3/papers.cfm?abstract_id=959023) [Accessed 18 May, 2024].
- Feldman, R. and Sanger, J., 2007. The Text Mining Handbook: Advanced Approaches in Analyzing Unstructured Data. Cambridge: Cambridge University Press. Online publication date: August 2009. DOI: 10.1017/CBO9780511546914 [Accessed 18 May, 2024].
- Tetlock, P.C., 2007. Giving content to investor sentiment: The role of media in the stock market. *The Journal of Finance*, 62(3), pp.1139-1168. First published: 08 May 2007. DOI: 10.1111/j.1540-6261.2007.01232.x [Accessed 18 May, 2024].
- Ngai, E.W.T., Hu, Y., Wong, Y.H., Chen, Y. and Sun, X., 2011. The application of data mining techniques in financial fraud detection: A classification framework and an academic review of literature. *Decision Support Systems*, 50(3), pp.559-569. Available at: <https://www.sciencedirect.com/science/article/pii/S0167923610001302> [Accessed 18 May, 2024].
- Goltz, N. and Mayo, M., 2017. Enhancing regulatory compliance by using artificial intelligence text mining to identify penalty clauses in legislation. *MIREL 2017 - Workshop on Mining and Reasoning with Legal texts*, [online] June 16th, 2017, London (UK). Available at: <https://ssrn.com/abstract=2977570> [Accessed 18 May, 2024].
- Bollen, J., Mao, H., & Zeng, X., 2011. Twitter mood predicts the stock market. *Journal of Computational Science*, 2(1), pp. 1-8. Available at: <https://doi.org/10.1016/j.jocs.2010.12.007> [Accessed 18 May 2024].
- Office of the Australian Information Commissioner (OAIC). (n.d.). Privacy Act 1988. Available at: <https://www.oaic.gov.au/privacy/privacy-legislation/the-privacy-act#:~:text=The%20Privacy%20Act%201988%20was,other%20organisations%2C%20handle%20personal%20information.> [Accessed 18 May 2024].
- Office of the Australian Information Commissioner (OAIC). (n.d.). Notifiable Data Breaches Scheme. Available at: <https://www.oaic.gov.au/privacy/notifiable-data-breaches> [Accessed 18 May 2024].
- Australian Government. (n.d.). Consumer Data Right. Available at: <https://www.cdr.gov.au/> [Accessed 18 May 2024].
- Australian Human Rights Commission. (2019). Human Rights and Technology. Available at:

<https://www.qld.gov.au/law/your-rights/human-rights#:~:text=The%20Human%20Rights%20Act%202019,protect%20and%20promote%20human%20rights> [Accessed 18 May 2024].

Australian Government. (2021). Artificial Intelligence Ethics Framework. Available at: <https://www.industry.gov.au/publications/australias-artificial-intelligence-ethics-framework> [Accessed 18 May 2024]. Date published: 7 November 2019.

Office of the Australian Information Commissioner (OAIC). (n.d.). Privacy resources. Available at: <https://www.oaic.gov.au/privacy> [Accessed 18 May 2024].

Australian Government. (2021). AI Ethics Principles. Available at: <https://www.industry.gov.au/publications/australias-artificial-intelligence-ethics-framework/australias-ai-ethics-principles> [Accessed 18 May 2024].

# Appendix

```
# ===== Task 3 =====

# Load the necessary libraries
library(ggplot2)
library(readr)
library(gridExtra)

# ===== 1. Load and Preprocess Data =====

# Load the dataset with suppressed column type message
data <- read_csv("UCI_credit_sample.csv", show_col_types = FALSE)

# Remove rows with any NA values
data <- na.omit(data)

# ===== 2. Create Individual Plots =====

# Box plot for LIMIT_BAL
plot1 <- ggplot(data, aes(x = as.factor(default), y = LIMIT_BAL, fill = as.factor(default))) +
  geom_boxplot() +
  labs(title = "Box Plot of LIMIT_BAL by Default", x = "Default", y = "LIMIT_BAL") +
  scale_fill_discrete(name = "Default")

# Box plot for AGE
plot2 <- ggplot(data, aes(x = as.factor(default), y = AGE, fill = as.factor(default))) +
  geom_boxplot() +
  labs(title = "Box Plot of AGE by Default", x = "Default", y = "AGE") +
  scale_fill_discrete(name = "Default")

# Box plot for BILL_AMT1
plot3 <- ggplot(data, aes(x = as.factor(default), y = BILL_AMT1, fill = as.factor(default))) +
  geom_boxplot() +
  labs(title = "Box Plot of BILL_AMT1 by Default", x = "Default", y = "BILL_AMT1") +
  scale_fill_discrete(name = "Default")

# Box plot for BILL_AMT2
plot4 <- ggplot(data, aes(x = as.factor(default), y = BILL_AMT2, fill = as.factor(default))) +
  geom_boxplot() +
  labs(title = "Box Plot of BILL_AMT2 by Default", x = "Default", y = "BILL_AMT2") +
  scale_fill_discrete(name = "Default")

# Box plot for PAY_AMT1
plot5 <- ggplot(data, aes(x = as.factor(default), y = PAY_AMT1, fill = as.factor(default))) +
  geom_boxplot() +
  labs(title = "Box Plot of PAY_AMT1 by Default", x = "Default", y = "PAY_AMT1") +
  scale_fill_discrete(name = "Default")

# Box plot for PAY_AMT2
plot6 <- ggplot(data, aes(x = as.factor(default), y = PAY_AMT2, fill = as.factor(default))) +
  geom_boxplot() +
  labs(title = "Box Plot of PAY_AMT2 by Default", x = "Default", y = "PAY_AMT2") +
  scale_fill_discrete(name = "Default")
```

```
# ===== 3. Arrange and Save Combined Box Plots =====
```

```
# Arrange plots in a 2x3 grid with adjusted height
```

```
combined_plot <- grid.arrange(plot1, plot2, plot3, plot4, plot5, plot6, ncol = 2, heights = c(1, 1, 1))
```

```
# Save the combined plot with increased height
```

```
ggsave("combined_plot.png", plot = combined_plot, width = 10, height = 12)
```

```
# ===== 4. Create Individual Bar Charts for Categorical Variables =====
```

```
# Bar chart for SEX
```

```
bar_chart_sex <- ggplot(data, aes(x = as.factor(SEX), fill = as.factor(default))) +  
  geom_bar(position = "dodge") +  
  labs(title = "Bar Chart of SEX by Default", x = "SEX", y = "Count") +  
  scale_fill_discrete(name = "Default")
```

```
# Bar chart for EDUCATION
```

```
bar_chart_education <- ggplot(data, aes(x = as.factor(EDUCATION), fill = as.factor(default))) +  
  geom_bar(position = "dodge") +  
  labs(title = "Bar Chart of EDUCATION by Default", x = "EDUCATION", y = "Count") +  
  scale_fill_discrete(name = "Default")
```

```
# Bar chart for MARRIAGE
```

```
bar_chart_marriage <- ggplot(data, aes(x = as.factor(MARRIAGE), fill = as.factor(default))) +  
  geom_bar(position = "dodge") +  
  labs(title = "Bar Chart of MARRIAGE by Default", x = "MARRIAGE", y = "Count") +  
  scale_fill_discrete(name = "Default")
```

```
# Bar chart for PAY_0
```

```
bar_chart_pay_0 <- ggplot(data, aes(x = PAY_0, fill = default)) +  
  geom_bar(position = "dodge") +  
  labs(title = "Bar Chart of PAY_0 by Default", x = "PAY_0", y = "Count") +  
  scale_fill_discrete(name = "Default")
```

```
# Bar chart for PAY_2
```

```
bar_chart_pay_2 <- ggplot(data, aes(x = PAY_2, fill = default)) +  
  geom_bar(position = "dodge") +  
  labs(title = "Bar Chart of PAY_2 by Default", x = "PAY_2", y = "Count") +  
  scale_fill_discrete(name = "Default")
```

```
# ===== 5. Arrange and Save Combined Bar Charts =====
```

```
# Arrange bar charts in a 2x3 grid
```

```
combined_bar_charts <- grid.arrange(  
  bar_chart_sex, bar_chart_education, bar_chart_marriage, bar_chart_pay_0, bar_chart_pay_2,  
  ncol = 3, nrow = 2  
)
```

```
# Save the combined bar charts with increased height
```

```
ggsave("combined_bar_charts.png", plot = combined_bar_charts, width = 10, height = 8)
```

```
# ===== 6. Generate Summary Statistics for Continuous Variables =====
```

```
summary_continuous <- data %>%
```

```
  summarise(  
    # Summary statistics for continuous variables  
    # Mean, standard deviation, minimum, maximum, and range  
    mean = mean(value),  
    sd = sd(value),  
    min = min(value),  
    max = max(value),  
    range = range(value)  
  )
```

```
LIMIT_BAL_min = min(LIMIT_BAL),
LIMIT_BAL_mean = mean(LIMIT_BAL),
LIMIT_BAL_median = median(LIMIT_BAL),
LIMIT_BAL_sd = sd(LIMIT_BAL),
LIMIT_BAL_max = max(LIMIT_BAL),
```

```
AGE_min = min(AGE),
AGE_mean = mean(AGE),
AGE_median = median(AGE),
AGE_sd = sd(AGE),
AGE_max = max(AGE),
```

```
BILL_AMT1_min = min(BILL_AMT1),
BILL_AMT1_mean = mean(BILL_AMT1),
BILL_AMT1_median = median(BILL_AMT1),
BILL_AMT1_sd = sd(BILL_AMT1),
BILL_AMT1_max = max(BILL_AMT1),
```

```
BILL_AMT2_min = min(BILL_AMT2),
BILL_AMT2_mean = mean(BILL_AMT2),
BILL_AMT2_median = median(BILL_AMT2),
BILL_AMT2_sd = sd(BILL_AMT2),
BILL_AMT2_max = max(BILL_AMT2),
```

```
PAY_AMT1_min = min(PAY_AMT1),
PAY_AMT1_mean = mean(PAY_AMT1),
PAY_AMT1_median = median(PAY_AMT1),
PAY_AMT1_sd = sd(PAY_AMT1),
PAY_AMT1_max = max(PAY_AMT1),
```

```
PAY_AMT2_min = min(PAY_AMT2),
PAY_AMT2_mean = mean(PAY_AMT2),
PAY_AMT2_median = median(PAY_AMT2),
PAY_AMT2_sd = sd(PAY_AMT2),
PAY_AMT2_max = max(PAY_AMT2)
```

```
)
```

# ===== 7. Generate Summary Statistics for Categorical Variables =====

```
summary_categorical <- data %>%
```

```
group_by(default) %>%
```

```
summarise(
```

```
SEX_1 = sum(SEX == 1),
```

```
SEX_2 = sum(SEX == 2),
```

```
EDUCATION_1 = sum(EDUCATION == 1),
```

```
EDUCATION_2 = sum(EDUCATION == 2),
```

```
EDUCATION_3 = sum(EDUCATION == 3),
```

```
EDUCATION_4 = sum(EDUCATION == 4),
```

```
EDUCATION_5 = sum(EDUCATION == 5),
```

```
EDUCATION_6 = sum(EDUCATION == 6),
```

```
MARRIAGE_0 = sum(MARRIAGE == 0),
```

```
MARRIAGE_1 = sum(MARRIAGE == 1),
```

```
MARRIAGE_2 = sum(MARRIAGE == 2),
```



```
MARRIAGE_3 = sum(MARRIAGE == 3),
```

```
PAY_0_0 = sum(PAY_0 == 0),
```

```
PAY_0_1 = sum(PAY_0 == 1),
```

```
PAY_0_2 = sum(PAY_0 == 2),
```

```
PAY_0_3 = sum(PAY_0 == 3),
```

```
PAY_0_4 = sum(PAY_0 == 4),
```

```
PAY_0_5 = sum(PAY_0 == 5),
```

```
PAY_0_6 = sum(PAY_0 == 6),
```

```
PAY_0_7 = sum(PAY_0 == 7),
```

```
PAY_0_8 = sum(PAY_0 == 8),
```

```
PAY_0_9 = sum(PAY_0 == 9),
```

```
PAY_2_0 = sum(PAY_2 == 0),
```

```
PAY_2_1 = sum(PAY_2 == 1),
```

```
PAY_2_2 = sum(PAY_2 == 2),
```

```
PAY_2_3 = sum(PAY_2 == 3),
```

```
PAY_2_4 = sum(PAY_2 == 4),
```

```
PAY_2_5 = sum(PAY_2 == 5),
```

```
PAY_2_6 = sum(PAY_2 == 6),
```

```
PAY_2_7 = sum(PAY_2 == 7),
```

```
PAY_2_8 = sum(PAY_2 == 8),
```

```
PAY_2_9 = sum(PAY_2 == 9)
```

```
)
```

```
# ===== Task 4 =====
```

```
# =====
```

```
# 1. Load Necessary Libraries
```

```
# =====
```

```
library(tidyverse)
```

```
library(TTR)
```

```
library(quantmod)
```

```
library(xts)
```

```
library(caret)
```

```
library(e1071)
```

```
library(nnet)
```

```
library(ggplot2)
```

```
# =====
```

```
# 2. Load and Preprocess the Data
```

```
# =====
```

```
# Load the data
```

```
data_stock_24 <- read.csv("data_stock_24.csv")
```

```
data_stock_24$Date <- as.Date(data_stock_24$Date, format="%Y-%m-%d")
```

```
# Inspect the column names to ensure correctness
```

```
colnames(data_stock_24)
```

```
# Subset the data to include only Date, ATVI-US, and sentiment columns
```

```
data_subset <- data_stock_24[, c("Date", "ATVI.US", "sentiment")]
```

```
# Convert data to xts format
```

```
data_xts <- xts(data_subset$`ATVI.US`, order.by=data_subset$Date)
```

```

# =====
# 3. Calculate Technical Indicators
# =====
# Calculate Daily Logarithmic Returns
log_returns <- ROC(data_xts, type="continuous", na.pad=FALSE)
colnames(log_returns) <- "log_returns"

# Calculate Moving Average (5-day)
ma_5 <- SMA(data_xts, n=5)
colnames(ma_5) <- "ma_5"

# Calculate Exponential Moving Average (5-day)
ema_5 <- EMA(data_xts, n=5)
colnames(ema_5) <- "ema_5"

# Calculate RSI (5-day)
rsi_5 <- RSI(data_xts, n=5)
colnames(rsi_5) <- "rsi_5"

# Combine all indicators into a single xts object
indicators <- merge(log_returns, ma_5, ema_5, rsi_5, all=FALSE)

# Lag the predictors by one period using xts::lag.xts, excluding log_returns
predictors <- indicators[, -1] # Exclude the first column (log_returns)
lagged_predictors <- xts::lag.xts(predictors, k=1)

# Combine the non-lagged log_returns with lagged predictors
final_indicators <- merge(indicators$log_returns, lagged_predictors, all=FALSE)

# Remove NA values resulting from the lag operation
final_indicators <- na.omit(final_indicators)

# =====
# 4. Process Sentiment Indicator
# =====
# Convert sentiment data to xts format
sentiment_xts <- xts(data_subset$sentiment, order.by=data_subset$Date)
colnames(sentiment_xts) <- "sentiment"

# Lag the sentiment indicator by one period
lagged_sentiment <- xts::lag.xts(sentiment_xts, k=1)

# Remove NA values resulting from the lag operation
lagged_sentiment <- na.omit(lagged_sentiment)

# =====
# 5. Merge All Indicators and Prepare Final Data
# =====
# Merge the lagged sentiment indicator with final indicators
final_data <- merge(final_indicators, lagged_sentiment, all=FALSE)

# Convert to data frame for easier handling
final_data_df <- data.frame(Date=index(final_data), coredata(final_data))

```

```

# =====
# 6. Set Seed and Split Data into Training and Testing Sets
# =====
# Set seed for reproducibility
set.seed(2345)

# Define the prediction window size
pred.window <- 100 # Last 100 days for the testing set

# Determine the indices for training and testing sets
idx <- c(1:(nrow(final_data_df) - pred.window))
d_train <- final_data_df[idx, ]
d_test <- final_data_df[(max(idx) + 1):nrow(final_data_df), ]

# Inspect the training and testing data
head(d_train)
head(d_test)

# =====
# 7. Neural Network Regression Forecasting
# =====
# Define the control parameters for model training
cntrl2 <- trainControl(method = "timeslice",
                      initialWindow = nrow(d_train) - pred.window,
                      horizon = pred.window,
                      fixedWindow = TRUE,
                      savePredictions = TRUE,
                      allowParallel = TRUE)

# Define preprocessing steps
prep1 <- c("center", "scale", "zv") # Include "zv" which removes variables with close to zero variance

# Train the neural network model
set.seed(2345)
nnet_model <- train(log_returns ~ . - Date,
                  data = d_train,
                  method = "nnet",
                  trControl = cntrl2,
                  preProcess = prep1,
                  linout = TRUE,
                  trace = FALSE,
                  metric = "RMSE",
                  tuneLength = 5)

# Print the trained model details
print(nnet_model)

# Prediction on the test set
nnet_predictions <- predict(nnet_model, newdata = d_test)

# Calculate performance measures
nnet_performance <- postResample(nnet_predictions, d_test$log_returns)

# Print performance measures

```

```

print(nnet_performance)

# Plot actual vs predicted values for Neural Network
data_pred_nnet <- data.frame(Date = d_test$Date, Actual = d_test$log_returns, Predicted = nnet_predictions)
data_pred_long_nnet <- pivot_longer(data_pred_nnet, cols = -Date)

plot_nn <- ggplot(data_pred_long_nnet, aes(x = Date, y = value, color = name)) +
  geom_line() +
  labs(title = "Neural Network Forecasting: Actual vs Predicted", x = "Date", y = "Log Returns") +
  theme_minimal()

# Save the plot using ggsave
ggsave("NN.png", plot = plot_nn, width = 10, height = 5)

# =====
# 8. SVM Regression Forecasting
# =====
# Train the SVM model with radial kernel
set.seed(2345)
svm_model <- train(log_returns ~ . - Date,
  data = d_train,
  method = "svmRadial",
  trControl = cntrl2,
  preProcess = prep1,
  metric = "RMSE",
  tuneLength = 5)

# Print the trained model details
print(svm_model)

# Prediction on the test set
svm_predictions <- predict(svm_model, newdata = d_test)

# Calculate performance measures
svm_performance <- postResample(svm_predictions, d_test$log_returns)

# Print performance measures
print(svm_performance)

# Plot actual vs predicted values for SVM
data_pred_svm <- data.frame(Date = d_test$Date, Actual = d_test$log_returns, Predicted = svm_predictions)
data_pred_long_svm <- pivot_longer(data_pred_svm, cols = -Date)

# Create the plot for SVM
plot_svm <- ggplot(data_pred_long_svm, aes(x = Date, y = value, color = name)) +
  geom_line() +
  labs(title = "SVM Regression Forecasting: Actual vs Predicted", x = "Date", y = "Log Returns") +
  theme_minimal()

# Save the plot using ggsave
ggsave("SVM.png", plot = plot_svm, width = 10, height = 5)

```

```
# Print performance measures for SVM
```

```
cat("SVM Performance Measures:\n")
```

```
print(svm_performance)
```

```
cat("Neural Network Performance Measures:\n")
```

```
print(nnet_performance)
```

